

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
```

```
In [2]: # Load the CSV file into a DataFrame
df = pd.read_csv('mcdonalds.csv')
```

```
In [3]: # Display the first few rows of the DataFrame to verify it's loaded correctly
print(df.head())
```

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	expensive	healthy	\
0	No	Yes	No	Yes	No	Yes	Yes	No	Yes	No	
1	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	Yes	No	
2	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes	Yes	
3	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	No	No	
4	No	Yes	No	Yes	Yes	Yes	Yes	No	No	Yes	

	disgusting	Like	Age	VisitFrequency	Gender
0	No	-3	61	Every three months	Female
1	No	+2	51	Every three months	Female
2	No	+1	62	Every three months	Female
3	Yes	+4	69	Once a week	Female
4	No	+2	49	Once a month	Male

```
In [4]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1453 entries, 0 to 1452
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   yummy                 1453 non-null  object
1   convenient            1453 non-null  object
2   spicy                 1453 non-null  object
3   fattening             1453 non-null  object
4   greasy                1453 non-null  object
5   fast                  1453 non-null  object
6   cheap                 1453 non-null  object
7   tasty                 1453 non-null  object
8   expensive             1453 non-null  object
9   healthy               1453 non-null  object
10  disgusting            1453 non-null  object
11  Like                  1453 non-null  object
12  Age                   1453 non-null  int64
13  VisitFrequency        1453 non-null  object
14  Gender                 1453 non-null  object
dtypes: int64(1), object(14)
memory usage: 170.4+ KB
None
```

```
In [5]: print(df.describe())
```

```

              Age
count  1453.000000
mean    44.604955
std     14.221178
min     18.000000
25%     33.000000
50%     45.000000
75%     57.000000
max     71.000000
```

```
In [6]: print(df.isnull().sum())
```

```

yummy          0
convenient     0
spicy          0
fattening      0
greasy         0
fast           0
cheap          0
tasty          0
expensive      0
healthy        0
disgusting     0
Like           0
Age            0
VisitFrequency 0
Gender         0
dtype: int64
```

```
In [7]: print(df.columns)
```

```
Index(['yummy', 'convenient', 'spicy', 'fattening', 'greasy', 'fast', 'cheap',
      'tasty', 'expensive', 'healthy', 'disgusting', 'Like', 'Age',
      'VisitFrequency', 'Gender'],
      dtype='object')
```

```
In [8]: print(df.dtypes)
```

```
yummy          object
convenient      object
spicy           object
fattening       object
greasy          object
fast            object
cheap           object
tasty           object
expensive       object
healthy         object
disgusting      object
Like            object
Age             int64
VisitFrequency  object
Gender          object
dtype: object
```

```
In [9]: # Check for duplicate rows
print(df.duplicated().sum())
```

```
22
```

```
In [10]: # Display unique values for each column
for column in df.columns:
    print(f"{column}: {df[column].nunique()} unique values")
```

```
yummy: 2 unique values
convenient: 2 unique values
spicy: 2 unique values
fattening: 2 unique values
greasy: 2 unique values
fast: 2 unique values
cheap: 2 unique values
tasty: 2 unique values
expensive: 2 unique values
healthy: 2 unique values
disgusting: 2 unique values
Like: 11 unique values
Age: 54 unique values
VisitFrequency: 6 unique values
Gender: 2 unique values
```

```
In [11]: # Preprocess the data (ensure non-numeric columns are handled appropriately, e.g., label encoding or dropping)
# For demonstration, assuming all columns are numeric or processed as such
data = df.select_dtypes(include=[np.number]) # Select only numeric columns for clustering
```

```
In [12]: # Standardize or normalize the data if necessary
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)
```

```
In [13]: # Run k-means for 2 to 8 segments with 10 random restarts (n_init)
segment_results = {}
for n_clusters in range(2, 9): # 2 to 8 clusters
    kmeans = KMeans(n_clusters=n_clusters, n_init=10, random_state=42)
    kmeans.fit(data_scaled)
    segment_results[n_clusters] = kmeans.labels_

# Print a summary for each k-means solution
print(f"Number of clusters: {n_clusters}")
print(f"Cluster labels: {np.unique(kmeans.labels_)}")
print()
```

Number of clusters: 2
Cluster labels: [0 1]

Number of clusters: 3
Cluster labels: [0 1 2]

Number of clusters: 4
Cluster labels: [0 1 2 3]

Number of clusters: 5
Cluster labels: [0 1 2 3 4]

Number of clusters: 6
Cluster labels: [0 1 2 3 4 5]

Number of clusters: 7
Cluster labels: [0 1 2 3 4 5 6]

Number of clusters: 8
Cluster labels: [0 1 2 3 4 5 6 7]

```
In [14]: # Optional: relabel segment numbers to be consistent across segmentations
# This step is complex and may require domain knowledge or matching centroids for consistency.

# Display the results for review
for n_clusters, labels in segment_results.items():
    df[f'Cluster_{n_clusters}'] = labels
```

```
In [15]: # Save or display DataFrame with added cluster columns for analysis
print(df.head())
```

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	expensive	healthy	\
0	No	Yes	No	Yes	No	Yes	Yes	No	Yes	No	
1	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	Yes	No	
2	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes	Yes	
3	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	No	No	
4	No	Yes	No	Yes	Yes	Yes	Yes	No	No	Yes	

	...	Age	VisitFrequency	Gender	Cluster_2	Cluster_3	Cluster_4	\
0	...	61	Every three months	Female	0	0	1	
1	...	51	Every three months	Female	0	2	2	
2	...	62	Every three months	Female	0	0	1	
3	...	69	Once a week	Female	0	0	1	
4	...	49	Once a month	Male	0	2	2	

	Cluster_5	Cluster_6	Cluster_7	Cluster_8
0	3	2	5	5
1	1	1	0	1
2	3	4	2	5
3	3	4	2	3
4	1	1	0	7

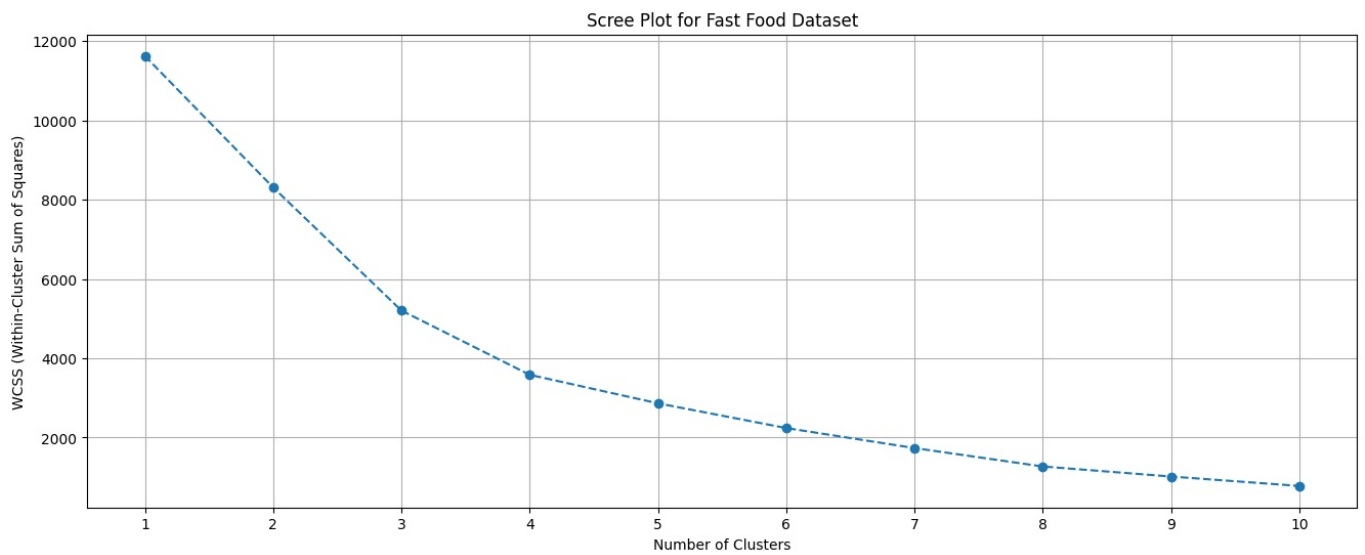
[5 rows x 22 columns]

```
In [16]: # Preprocess the data (select numeric columns)
data = df.select_dtypes(include=[np.number])
```

```
In [17]: # Standardize the data
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)
```

```
In [18]: # Calculate WCSS for different number of clusters
wcss = []
for i in range(1, 11): # Range from 1 to 10 clusters
    kmeans = KMeans(n_clusters=i, n_init=10, random_state=42)
    kmeans.fit(data_scaled)
    wcss.append(kmeans.inertia_)
```

```
In [19]: # Plot the scree plot
plt.figure(figsize=(16, 6))
plt.plot(range(1, 11), wcss, marker='o', linestyle='--')
plt.title('Scree Plot for Fast Food Dataset')
plt.xlabel('Number of Clusters')
plt.ylabel('WCSS (Within-Cluster Sum of Squares)')
plt.xticks(range(1, 11))
plt.grid(True)
plt.show()
```



```
In [20]: from sklearn.metrics import adjusted_rand_score
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Preprocess the data (select numeric columns)
data = df.select_dtypes(include=[np.number])
```

```
In [21]: # Standardize the data
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)
```

```
In [22]: # Run k-means clustering multiple times and evaluate stability
num_clusters = 4 # Choose the number of clusters to evaluate (can be adjusted)
num_runs = 10    # Number of runs for evaluating stability
cluster_labels_list = []
```

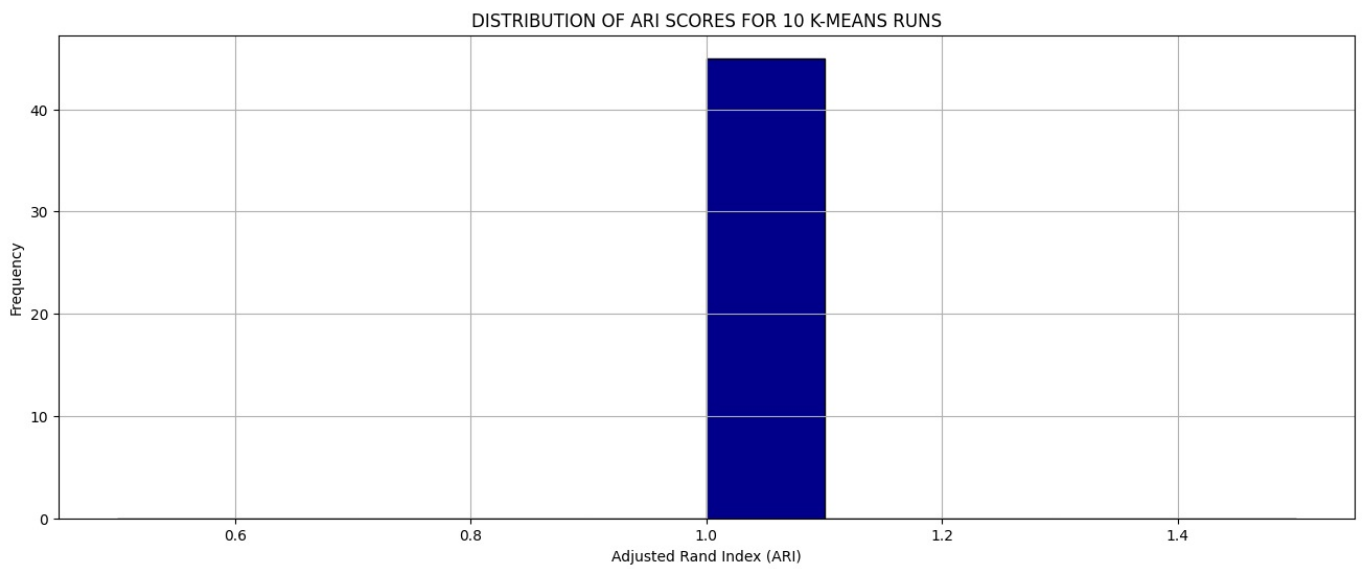
```
In [23]: # Run k-means multiple times
for i in range(num_runs):
    kmeans = KMeans(n_clusters=num_clusters, n_init=10, random_state=i)
    kmeans.fit(data_scaled)
    cluster_labels_list.append(kmeans.labels_)
```

```
In [24]: # Compute the Adjusted Rand Index (ARI) between all pairs of clustering results
ari_scores = []
for i in range(len(cluster_labels_list)):
    for j in range(i + 1, len(cluster_labels_list)):
        ari = adjusted_rand_score(cluster_labels_list[i], cluster_labels_list[j])
        ari_scores.append(ari)
```

```
In [25]: # Display the average ARI score
average_ari = np.mean(ari_scores)
print(f"Average Adjusted Rand Index (ARI) for {num_runs} runs with {num_clusters} clusters: {average_ari:.3f}")
```

Average Adjusted Rand Index (ARI) for 10 runs with 4 clusters: 1.000

```
In [26]: # Plot the distribution of ARI scores
plt.figure(figsize=(16, 6))
plt.hist(ari_scores, bins=10, color='darkblue', edgecolor='black')
plt.title(f'DISTRIBUTION OF ARI SCORES FOR {num_runs} K-MEANS RUNS')
plt.xlabel('Adjusted Rand Index (ARI)')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()
```



```
In [27]: import seaborn as sns
from sklearn.preprocessing import StandardScaler

# Preprocess the data (select numeric columns)
data = df.select_dtypes(include=[np.number])
```

```
In [28]: # Standardize the data
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)
```

```
In [29]: # Run k-means clustering for 4 segments
num_clusters = 4
kmeans = KMeans(n_clusters=num_clusters, n_init=10, random_state=42)
df['Cluster'] = kmeans.fit_predict(data_scaled)
```

```
In [30]: # Create a gorge plot for the four-segment k-means solution
plt.figure(figsize=(15, 12))
for feature in data.columns:
    plt.subplot(len(data.columns) // 2 + 1, 2, list(data.columns).index(feature) + 1)
    sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
    plt.title(f'Distribution of {feature} by Cluster')
    plt.xlabel('Cluster')
    plt.ylabel(feature)

plt.tight_layout()
plt.suptitle('Gorge Plot of Four-Segment K-Means Solution', y=1.02, fontsize=16)
plt.show()
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

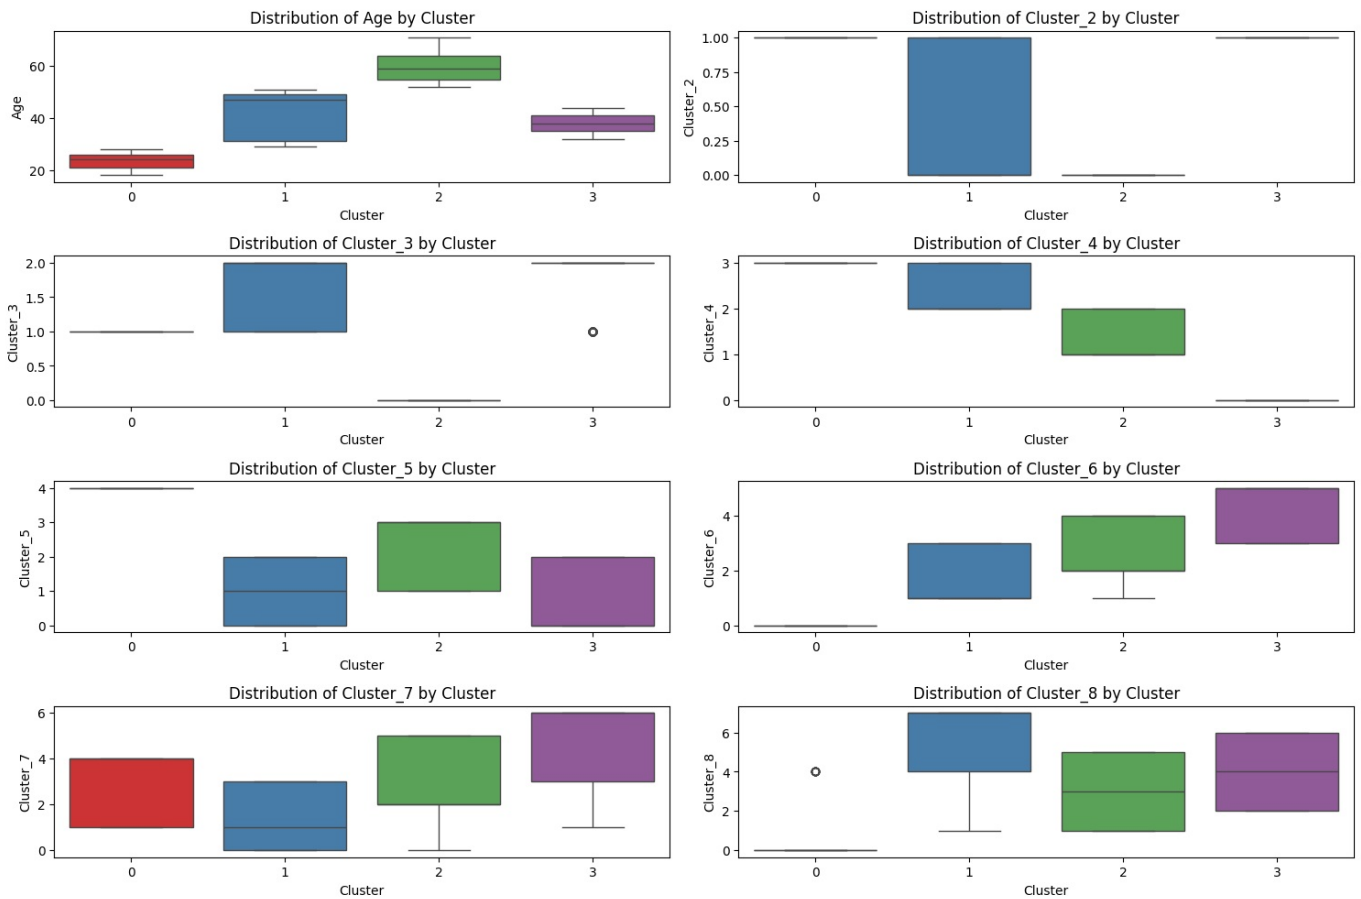
```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

C:\Users\LOKESH\AppData\Local\Temp\ipykernel_47236\284269802.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cluster', y=feature, data=df, palette='Set1')
```

Gorge Plot of Four-Segment K-Means Solution



```
In [31]: from sklearn.metrics import adjusted_rand_score
from sklearn.preprocessing import StandardScaler

# Preprocess the data (select numeric columns)
data = df.select_dtypes(include=[np.number])

# Standardize the data
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)

# Run k-means for solutions from 2 to 8 segments and store the labels
solutions = {}
for n_clusters in range(2, 9):
    kmeans = KMeans(n_clusters=n_clusters, n_init=10, random_state=42)
    kmeans.fit(data_scaled)
    solutions[n_clusters] = kmeans.labels_
```

```
In [32]: # Calculate segment level stability
segment_stability = {}
for current_clusters in range(3, 8):
    previous_labels = solutions[current_clusters - 1]
    current_labels = solutions[current_clusters]
    next_labels = solutions[current_clusters + 1]

    stability_scores = []
    for cluster in np.unique(current_labels):
        # Find the indices of data points in the current cluster
        indices = np.where(current_labels == cluster)[0]

        # Calculate ARI for current cluster against previous and next solutions
        ari_prev = adjusted_rand_score(previous_labels[indices], current_labels[indices])
        ari_next = adjusted_rand_score(current_labels[indices], next_labels[indices])

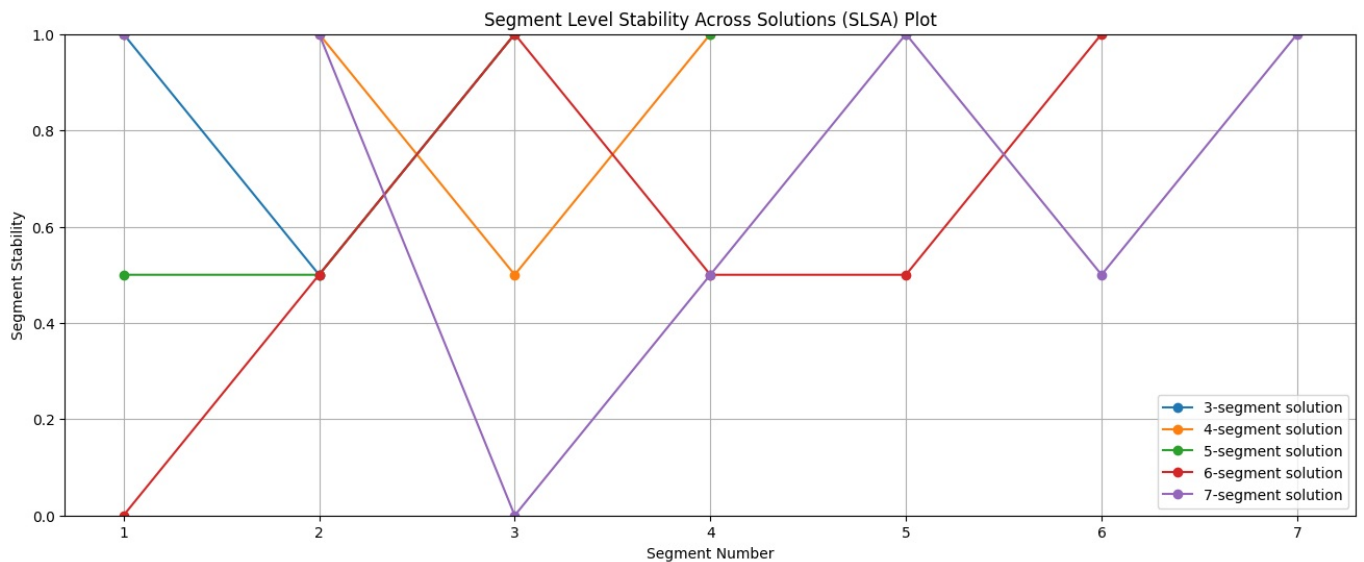
        # Average stability for the current segment
        stability_score = (ari_prev + ari_next) / 2
        stability_scores.append(stability_score)

    segment_stability[current_clusters] = stability_scores
```

```
In [33]: # Plot the SLSA plot
plt.figure(figsize=(16, 6))
for clusters, stability in segment_stability.items():
    plt.plot(range(1, len(stability) + 1), stability, marker='o', label=f'{clusters}-segment solution')

plt.ylim(0, 1)
plt.xlabel("Segment Number")
```

```
plt.ylabel("Segment Stability")
plt.title("Segment Level Stability Across Solutions (SLSA) Plot")
plt.legend()
plt.grid(True)
plt.show()
```

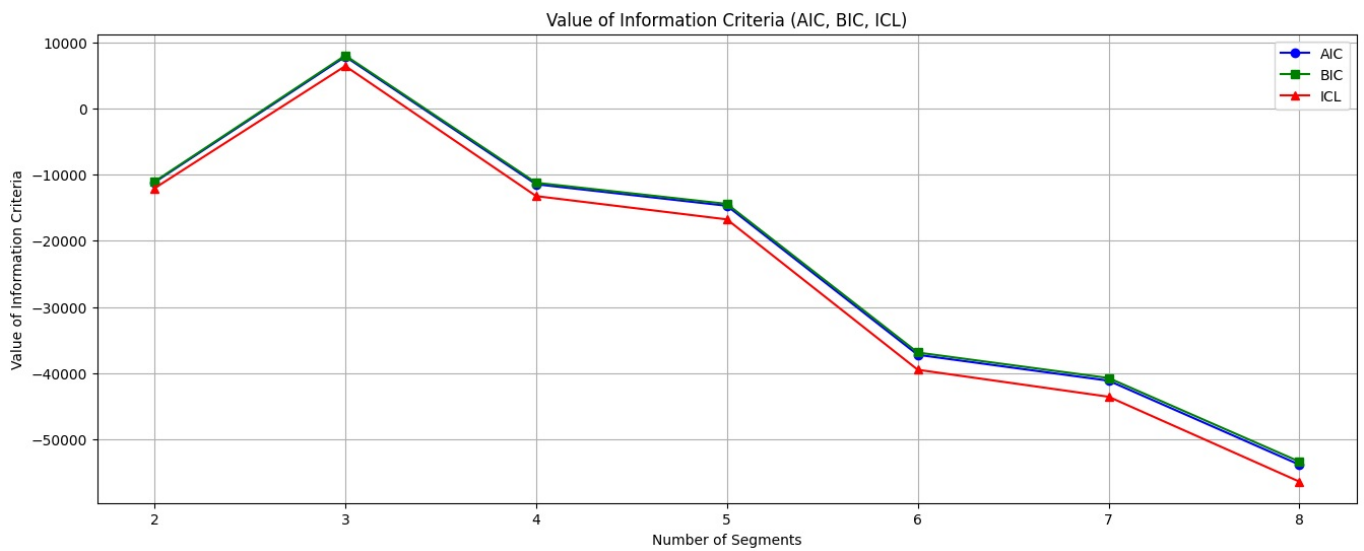


```
In [34]: from sklearn.mixture import GaussianMixture
import matplotlib.pyplot as plt
from sklearn.preprocessing import Binarizer, StandardScaler

# Preprocess the data (select numeric columns and binarize them)
data = df.select_dtypes(include=[np.number])
binarizer = Binarizer(threshold=0.5) # Adjust threshold as necessary
data_binarized = binarizer.fit_transform(StandardScaler().fit_transform(data))
# Fit mixture models for 2 to 8 segments
aic_values = []
bic_values = []
icl_values = [] # In Python, ICL may be approximated by custom calculations
```

```
In [35]: for n_components in range(2, 9):
    gmm = GaussianMixture(n_components=n_components, covariance_type='spherical', random_state=1234, n_init=10)
    gmm.fit(data_binarized)
    aic_values.append(gmm.aic(data_binarized))
    bic_values.append(gmm.bic(data_binarized))
    # ICL can be approximated using penalized BIC; here is a simplified approach
    icl_values.append(gmm.bic(data_binarized) - np.log(n_components) * len(data_binarized))
```

```
In [36]: # Plot the information criteria
plt.figure(figsize=(16, 6))
plt.plot(range(2, 9), aic_values, marker='o', label='AIC', color='blue')
plt.plot(range(2, 9), bic_values, marker='s', label='BIC', color='green')
plt.plot(range(2, 9), icl_values, marker='^', label='ICL', color='red')
plt.title('Value of Information Criteria (AIC, BIC, ICL)')
plt.xlabel('Number of Segments')
plt.ylabel('Value of Information Criteria')
plt.xticks(range(2, 9))
plt.legend()
plt.grid(True)
plt.show()
```

In [37]: # Preprocess the data (select numeric columns and binarize them)

```
data = df.select_dtypes(include=[np.number])
scaler = StandardScaler()
data_scaled = scaler.fit_transform(data)
binarizer = Binarizer(threshold=0.5)
data_binarized = binarizer.fit_transform(data_scaled)

# Initialize lists to store AIC, BIC, and ICL values
aic_values = []
bic_values = []
icl_values = [] # ICL approximation using a penalized BIC
```

In [38]: # Check the first few rows of the dataset

```
print(df.head())

# Ensure the 'Like' column exists
if 'Like' not in df.columns:
    raise ValueError("The 'Like' column is missing in the dataset.")

# Convert the dependent variable 'LIKE' to a numeric variable
# Mapping the categorical variable with 11 levels from 'ILOVEIT!' (+5) to 'IHATE IT!' (-5)
like_mapping = {
    'I LOVE IT!': 5, '4': 4, '3': 3, '2': 2, '1': 1, '0': 0,
    '-1': -1, '-2': -2, '-3': -3, '-4': -4, 'I HATE IT!': -5
}
```

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	expensive	healthy	\
0	No	Yes	No	Yes	No	Yes	Yes	No	Yes	No	
1	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	Yes	No	
2	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes	Yes	
3	Yes	Yes	No	Yes	Yes	Yes	Yes	Yes	No	No	
4	No	Yes	No	Yes	Yes	Yes	Yes	No	No	Yes	

	...	VisitFrequency	Gender	Cluster_2	Cluster_3	Cluster_4	Cluster_5	\
0	...	Every three months	Female	0	0	1	3	
1	...	Every three months	Female	0	2	2	1	
2	...	Every three months	Female	0	0	1	3	
3	...	Once a week	Female	0	0	1	3	
4	...	Once a month	Male	0	2	2	1	

	Cluster_6	Cluster_7	Cluster_8	Cluster
0	2	5	5	2
1	1	0	1	1
2	4	2	5	2
3	4	2	3	2
4	1	0	7	1

[5 rows x 23 columns]

In [39]: # Replace the categorical 'Like' with numeric 'Like.n'

```
df['Like.n'] = df['Like'].map(like_mapping)

# Check the mapping results
print(df['Like.n'].value_counts().sort_index())

# Define the independent variables (perceptions of McDonald's)
independent_vars = ['yummy', 'convenient', 'spicy', 'fattening', 'greasy',
                    'fast', 'cheap', 'tasty', 'expensive', 'healthy', 'disgusting']

# Check if all independent variables exist in the DataFrame
missing_vars = [var for var in independent_vars if var not in df.columns]
```

```
if missing_vars:
    raise ValueError(f"The following independent variables are missing: {missing_vars}")
```

```
Like.n
-4.0    71
-3.0    73
-2.0    59
-1.0    58
 0.0   169
Name: count, dtype: int64
```

```
In [40]: import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# Example data generation (replace this with your actual data)
np.random.seed(0)
X1 = np.random.rand(100, 1) * 10
y1 = 2 * X1 + 5 + np.random.randn(100, 1) # Segment 1
X2 = np.random.rand(100, 1) * 10
y2 = -3 * X2 + 30 + np.random.randn(100, 1) # Segment 2

X = np.vstack((X1, X2))
y = np.vstack((y1, y2))
```

```
In [41]: # Create a function to fit the two segments
def fit_mixture_regression(X, y):
    # Fit the first linear regression model
    model1 = LinearRegression()
    model1.fit(X[:100], y[:100])

    # Fit the second linear regression model
    model2 = LinearRegression()
    model2.fit(X[100:], y[100:])


    return model1, model2
```

```
In [42]: # Fit the mixture regression
model1, model2 = fit_mixture_regression(X, y)

# Print the regression coefficients
print("Segment 1 Coefficients (Intercept, Slope):", model1.intercept_[0], model1.coef_[0][0])
print("Segment 2 Coefficients (Intercept, Slope):", model2.intercept_[0], model2.coef_[0][0])

# Plotting the results
plt.scatter(X, y, color='blue', label='Data points')
plt.plot(X[:100], model1.predict(X[:100]), color='red', label='Segment 1 Fit')
plt.plot(X[100:], model2.predict(X[100:]), color='green', label='Segment 2 Fit')
plt.xlabel('X')
plt.ylabel('y')
plt.title('Two-Segment Mixture of Linear Regression')
plt.legend()
plt.show()
```

```
Segment 1 Coefficients (Intercept, Slope): 5.222151077447226 1.9936935021402045
Segment 2 Coefficients (Intercept, Slope): 30.101897754959293 -3.044003252201276
```



```
In [43]: import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

```
In [44]: # Example data generation (replace this with your actual data)
np.random.seed(0)
X1 = np.random.rand(100, 1) * 10
y1 = 2 * X1 + 5 + np.random.randn(100, 1) # Segment 1
X2 = np.random.rand(100, 1) * 10
y2 = -3 * X2 + 30 + np.random.randn(100, 1) # Segment 2

X = np.vstack((X1, X2))
y = np.vstack((y1, y2))
```

```
In [45]: # Create a function to fit the two segments
def fit_mixture_regression(X, y):
    # Fit the first linear regression model
    model1 = LinearRegression()
    model1.fit(X[:100], y[:100])

    # Fit the second linear regression model
    model2 = LinearRegression()
    model2.fit(X[100:], y[100:])

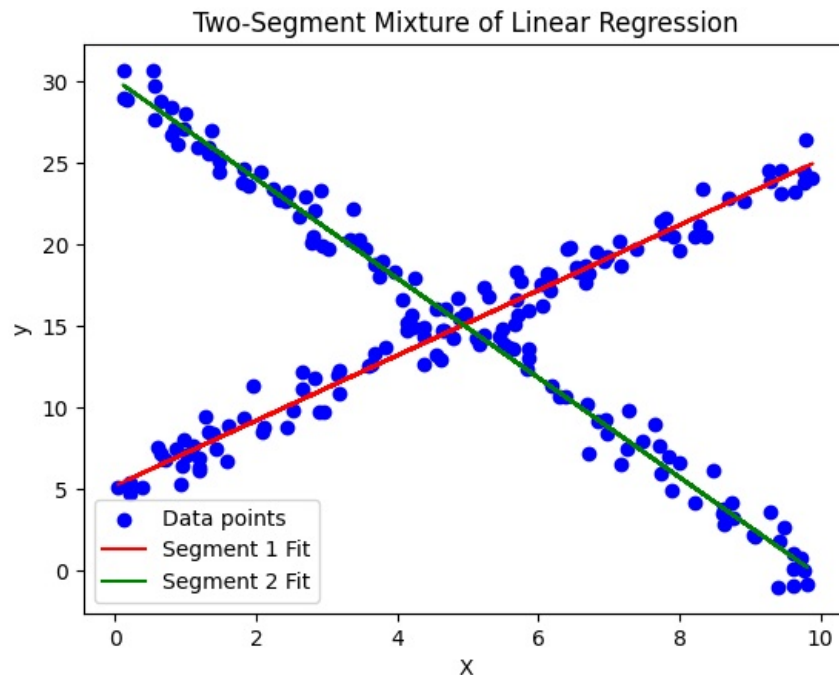
    return model1, model2

# Fit the mixture regression
model1, model2 = fit_mixture_regression(X, y)
```

```
In [46]: # Print the regression coefficients
print("Segment 1 Coefficients (Intercept, Slope):", model1.intercept_[0], model1.coef_[0][0])
print("Segment 2 Coefficients (Intercept, Slope):", model2.intercept_[0], model2.coef_[0][0])

# Plotting the results
plt.scatter(X, y, color='blue', label='Data points')
plt.plot(X[:100], model1.predict(X[:100]), color='red', label='Segment 1 Fit')
plt.plot(X[100:], model2.predict(X[100:]), color='green', label='Segment 2 Fit')
plt.xlabel('X')
plt.ylabel('y')
plt.title('Two-Segment Mixture of Linear Regression')
plt.legend()
plt.show()
```

Segment 1 Coefficients (Intercept, Slope): 5.222151077447226 1.9936935021402045
Segment 2 Coefficients (Intercept, Slope): 30.101897754959293 -3.044003252201276



```
In [47]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
```

```
In [48]: # Example data generation (replace this with your actual data)
np.random.seed(0)
data = {
    'disgusting': np.random.rand(100),
    'expensive': np.random.rand(100),
}
```

```

'greasy': np.random.rand(100),
'healthy': np.random.rand(100),
'spicy': np.random.rand(100),
'fast': np.random.rand(100),
'convenient': np.random.rand(100),
'fattening': np.random.rand(100),
'cheap': np.random.rand(100),
'tasty': np.random.rand(100),
'yummy': np.random.rand(100)
}
df = pd.DataFrame(data)

```

```

In [49]: # Perform K-Means clustering
kmeans = KMeans(n_clusters=4, random_state=0)
df['segment'] = kmeans.fit_predict(df)

# Calculate mean values for each segment
segment_profile = df.groupby('segment').mean().reset_index()


# Plot the segment profile
plt.figure(figsize=(14, 8))
for i, segment in segment_profile.iterrows():
    plt.plot(segment_profile.columns[1:], segment.values[1:], marker='o', label=f'Segment {int(segment["segment"])}')

plt.title('Segment Profile Plot for the Four-Segment Solution')
plt.xlabel('Attributes')
plt.ylabel('Mean Values')
plt.xticks(rotation=45)
plt.legend(title='Segments')
plt.grid(True)
plt.tight_layout()
plt.show()

```

C:\Users\LOKESH\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning

```
super()._check_params_vs_input(X, default_n_init=10)
```



```

In [50]: import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

# Example data generation (replace this with your actual data)
np.random.seed(0)
data = {
    'disgusting': np.random.rand(100),
    'expensive': np.random.rand(100),
    'greasy': np.random.rand(100),
    'healthy': np.random.rand(100),
    'spicy': np.random.rand(100),
    'fast': np.random.rand(100),
    'convenient': np.random.rand(100),
    'fattening': np.random.rand(100),

```

```

'cheap': np.random.rand(100),
'tasty': np.random.rand(100),
'yummy': np.random.rand(100)
}
df = pd.DataFrame(data)

```

```

In [51]: # Perform K-Means clustering
kmeans = KMeans(n_clusters=4, random_state=0)
df['segment'] = kmeans.fit_predict(df)

# Perform PCA for dimensionality reduction
pca = PCA(n_components=2)
pca_result = pca.fit_transform(df.drop('segment', axis=1))

# Add PCA results to the DataFrame
df['PCA1'] = pca_result[:, 0]
df['PCA2'] = pca_result[:, 1]

# Plot the segments in the PCA space
plt.figure(figsize=(10, 7))
colors = ['red', 'green', 'blue', 'orange']

```

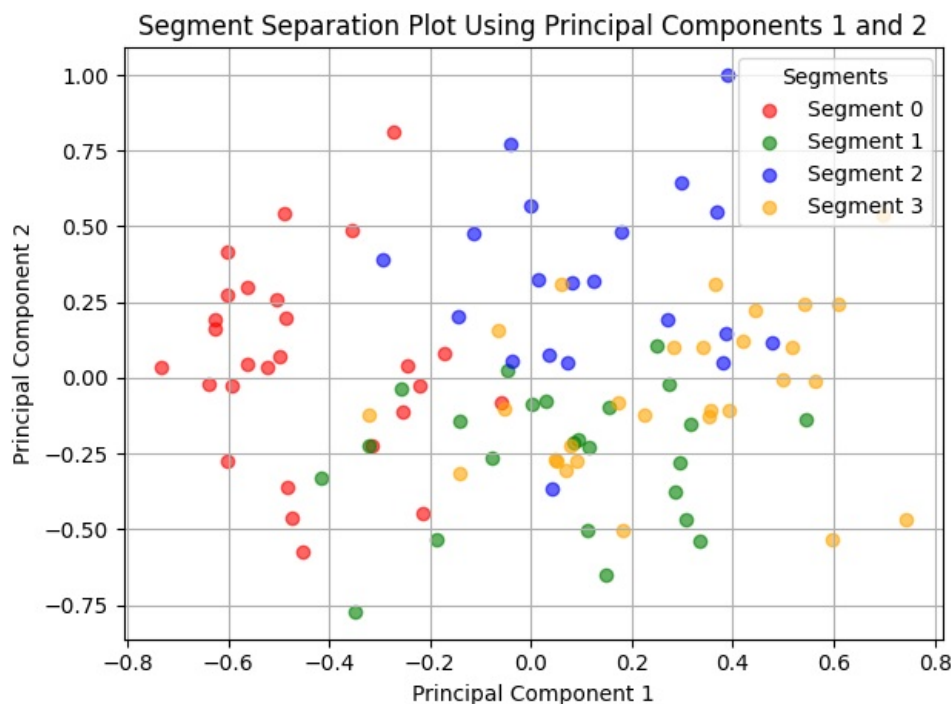
C:\Users\LOKESH\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super().check_params_vs_input(X, default_n_init=10)
<Figure size 1000x700 with 0 Axes>

```

In [52]: for i in range(4): # Number of clusters/segments
plt.scatter(
    df[df['segment'] == i]['PCA1'],
    df[df['segment'] == i]['PCA2'],
    label=f'Segment {i}',
    color=colors[i],
    alpha=0.6
)

plt.title('Segment Separation Plot Using Principal Components 1 and 2')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend(title='Segments')
plt.grid(True)
plt.tight_layout()
plt.show()

```



```

In [53]: import matplotlib.pyplot as plt
from statsmodels.graphics.mosaicplot import mosaic
from sklearn.cluster import KMeans

# Example data generation (replace this with your actual data)
np.random.seed(0)
data = {
    'ILIKEIT': np.random.choice(['ILOVEIT!', 'Somewhat Like', 'Neutral', 'Dislike', 'IHATEIT!'], size=100),
    'disgusting': np.random.rand(100),
    'expensive': np.random.rand(100),

```

```

'greasy': np.random.rand(100),
'healthy': np.random.rand(100),
'spicy': np.random.rand(100),
'fast': np.random.rand(100),
'convenient': np.random.rand(100),
'fattening': np.random.rand(100),
'cheap': np.random.rand(100),
'tasty': np.random.rand(100),
'yummy': np.random.rand(100)
}
df = pd.DataFrame(data)

```

```


In [54]: # Perform K-Means clustering
kmeans = KMeans(n_clusters=4, random_state=0)
df['segment'] = kmeans.fit_predict(df.drop('ILIKEIT', axis=1))

# Create a cross-tabulation and convert to a dictionary
crosstab = pd.crosstab(df['segment'], df['ILIKEIT'])
crosstab_dict = crosstab.stack().to_dict()

# Plot the mosaic plot
plt.figure(figsize=(12, 8))
mosaic(crosstab_dict, title='Shaded Mosaic Plot for Segment Membership and ILIKEIT', gap=0.02)
plt.show()

```

C:\Users\LOKESH\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super().check_params_vs_input(X, default_n_init=10)
<Figure size 1200x800 with 0 Axes>



```

In [55]: import matplotlib.pyplot as plt
from statsmodels.graphics.mosaicplot import mosaic
from sklearn.cluster import KMeans

# Example data generation (replace this with your actual data)
np.random.seed(0)
data = {
    'Gender': np.random.choice(['Male', 'Female'], size=100),
    'disgusting': np.random.rand(100),
    'expensive': np.random.rand(100),
    'greasy': np.random.rand(100),
    'healthy': np.random.rand(100),
    'spicy': np.random.rand(100),
    'fast': np.random.rand(100),
    'convenient': np.random.rand(100),
    'fattening': np.random.rand(100),
    'cheap': np.random.rand(100),
    'tasty': np.random.rand(100),
    'yummy': np.random.rand(100)
}
df = pd.DataFrame(data)

```

```

In [56]: # Perform K-Means clustering
kmeans = KMeans(n_clusters=4, random_state=0)
df['segment'] = kmeans.fit_predict(df.drop('Gender', axis=1))


```

```
# Create a cross-tabulation and convert to a dictionary
crosstab = pd.crosstab(df['segment'], df['Gender'])
crosstab_dict = crosstab.stack().to_dict()

# Plot the mosaic plot
plt.figure(figsize=(12, 8))
mosaic(crosstab_dict, title='Shaded Mosaic Plot for Segment Membership and Gender', gap=0.02)
plt.show()
```

C:\Users\LOKESH\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

```
super()._check_params_vs_input(X, default_n_init=10)
<Figure size 1200x800 with 0 Axes>
```



```
In [57]: import matplotlib.pyplot as plt
from sklearn.cluster import KMeans


# Example data generation (replace this with your actual data)
np.random.seed(0)
data = {
    'Age': np.random.randint(18, 60, size=100),
    'disgusting': np.random.rand(100),
    'expensive': np.random.rand(100),
    'greasy': np.random.rand(100),
    'healthy': np.random.rand(100),
    'spicy': np.random.rand(100),
    'fast': np.random.rand(100),
    'convenient': np.random.rand(100),
    'fattening': np.random.rand(100),
    'cheap': np.random.rand(100),
    'tasty': np.random.rand(100),
    'yummy': np.random.rand(100)
}
df = pd.DataFrame(data)
```

```
In [58]: # Perform K-Means clustering
kmeans = KMeans(n_clusters=4, random_state=0)
df['segment'] = kmeans.fit_predict(df.drop('Age', axis=1))

# Create a parallel box-and-whisker plot for age by segment
plt.figure(figsize=(12, 8))
df.boxplot(column='Age', by='segment', grid=False)
plt.title('Parallel Box-and-Whisker Plot of Age by Segment')
plt.suptitle('') # Remove the default subtitle
plt.xlabel('Segment')
plt.ylabel('Age')
plt.show()
```

C:\Users\LOKESH\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

```
super()._check_params_vs_input(X, default_n_init=10)
<Figure size 1200x800 with 0 Axes>
```

```
In [59]: import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

# Example data generation
np.random.seed(0)
data = {
    'Age': np.random.randint(18, 60, size=100),
    'disgusting': np.random.rand(100),
    'expensive': np.random.rand(100),
    'greasy': np.random.rand(100),
    'healthy': np.random.rand(100),
    'spicy': np.random.rand(100),
    'fast': np.random.rand(100),
    'convenient': np.random.rand(100),
    'fattening': np.random.rand(100),
    'cheap': np.random.rand(100),
    'tasty': np.random.rand(100),
    'yummy': np.random.rand(100),
}
df = pd.DataFrame(data)
```

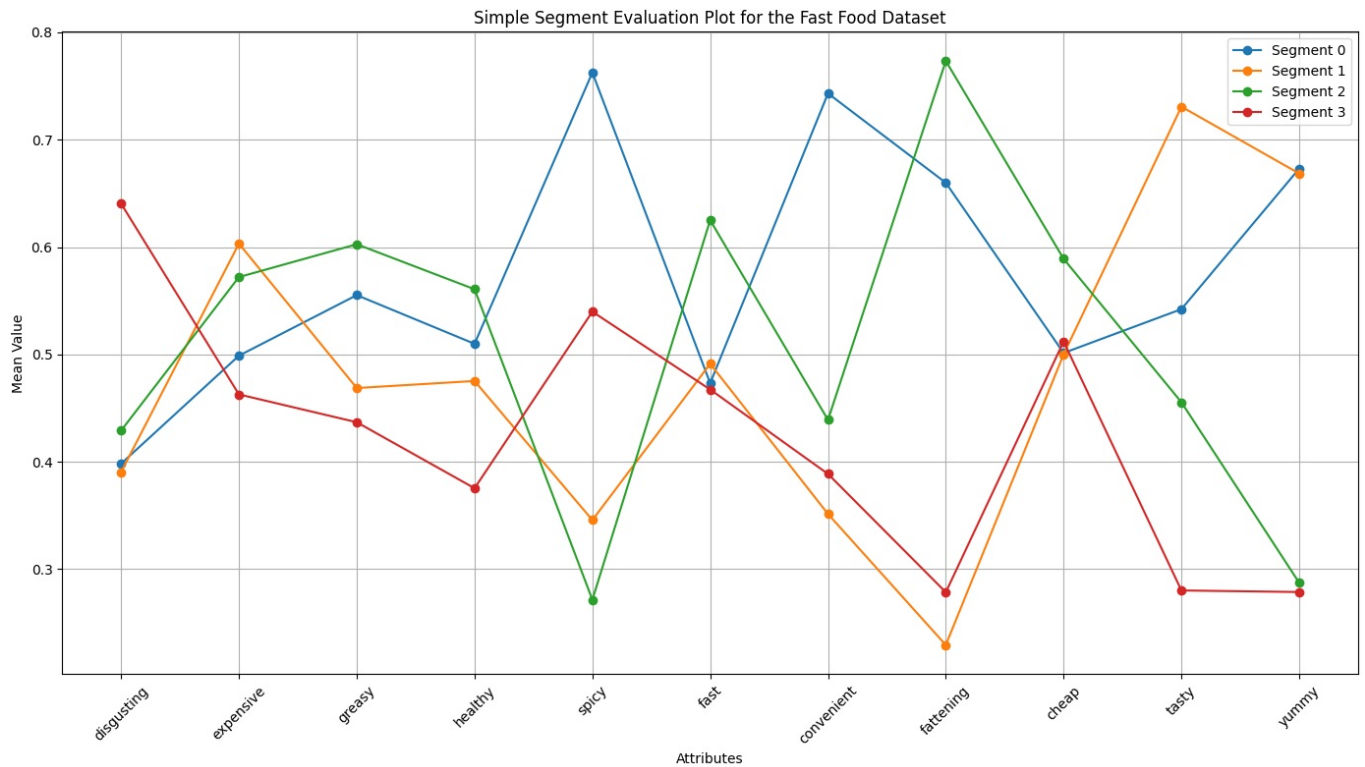
```
In [60]: import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
np.random.seed(0)
data = {
    'disgusting': np.random.rand(100),
    'expensive': np.random.rand(100),
    'greasy': np.random.rand(100),
    'healthy': np.random.rand(100),
    'spicy': np.random.rand(100),
    'fast': np.random.rand(100),
    'convenient': np.random.rand(100),
    'fattening': np.random.rand(100),
    'cheap': np.random.rand(100),
    'tasty': np.random.rand(100),
    'yummy': np.random.rand(100),
}
df = pd.DataFrame(data)
```

```
In [61]: # Perform K-Means clustering to create segments
kmeans = KMeans(n_clusters=4, random_state=0)
df['segment'] = kmeans.fit_predict(df)
# Calculate the mean of each feature for each segment
segment_profile = df.groupby('segment').mean()
plt.figure(figsize=(14, 8))
for segment in segment_profile.index:
    plt.plot(segment_profile.columns, segment_profile.loc[segment], marker='o', label=f'Segment {segment}')
plt.title('Simple Segment Evaluation Plot for the Fast Food Dataset')
plt.xlabel('Attributes')
```



```
plt.ylabel('Mean Value')
plt.xticks(rotation=45)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

C:\Users\LOKESH\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:1412: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)



Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js